

# REFORMER :

# The efficient Transformer

Blog: <https://ai.googleblog.com/2020/01/reformer-efficient-transformer.html>

Paper: <https://arxiv.org/abs/2001.04451>

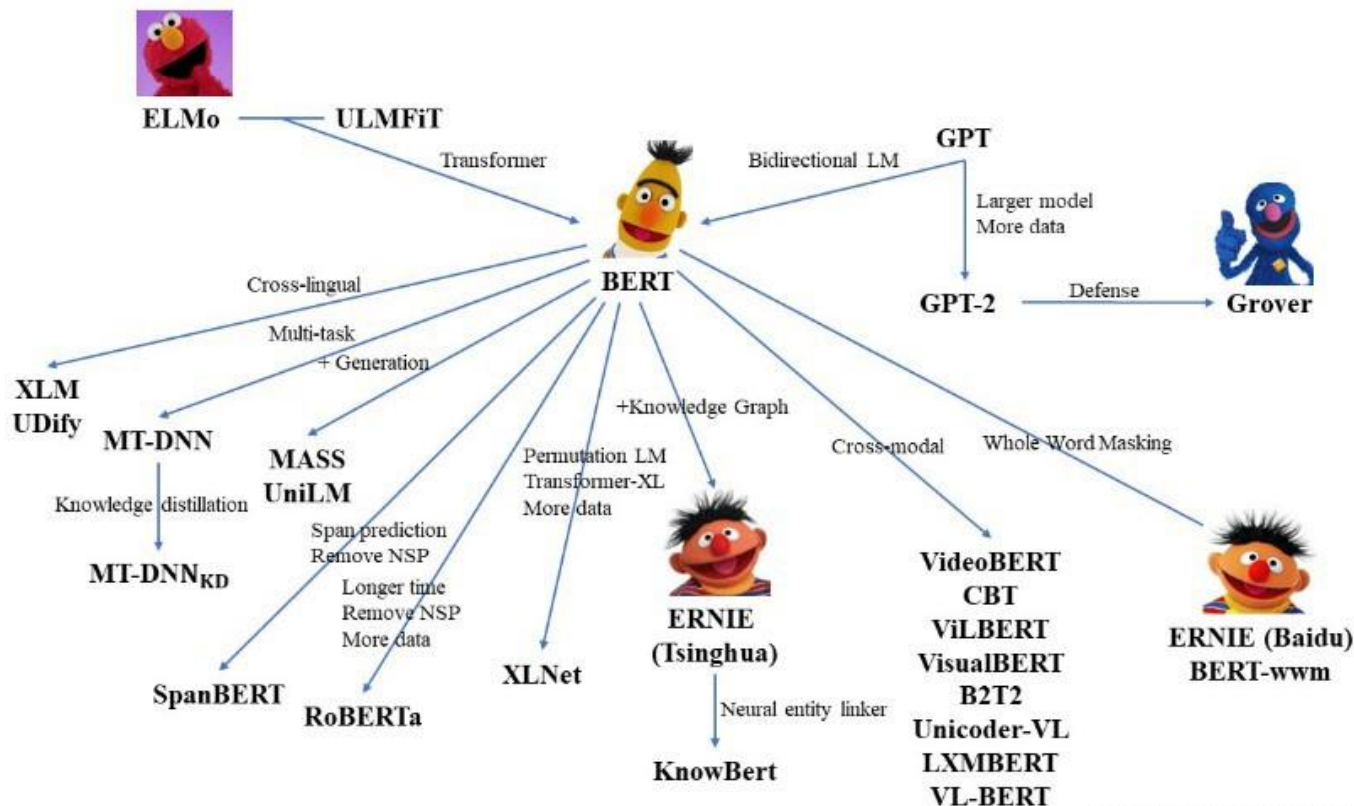
Openreview: <https://openreview.net/forum?id=rkgNKkHtvB>

Oral paper at ICLR 2020, spotlight

# Slide index

- 1) Some background on state of the art NLP and the issues : quadratic memory growth
- 2) Reduce memory usage using hash and divide-and-conquer
- 3) Reduce more memory by implementing reversible transformer
- 4) Ablation study results
- 5) Summary
- 6) Reference

# The **Cool kids** in Natural Language Processing



# The **Cool kids** in Natural Language Processing

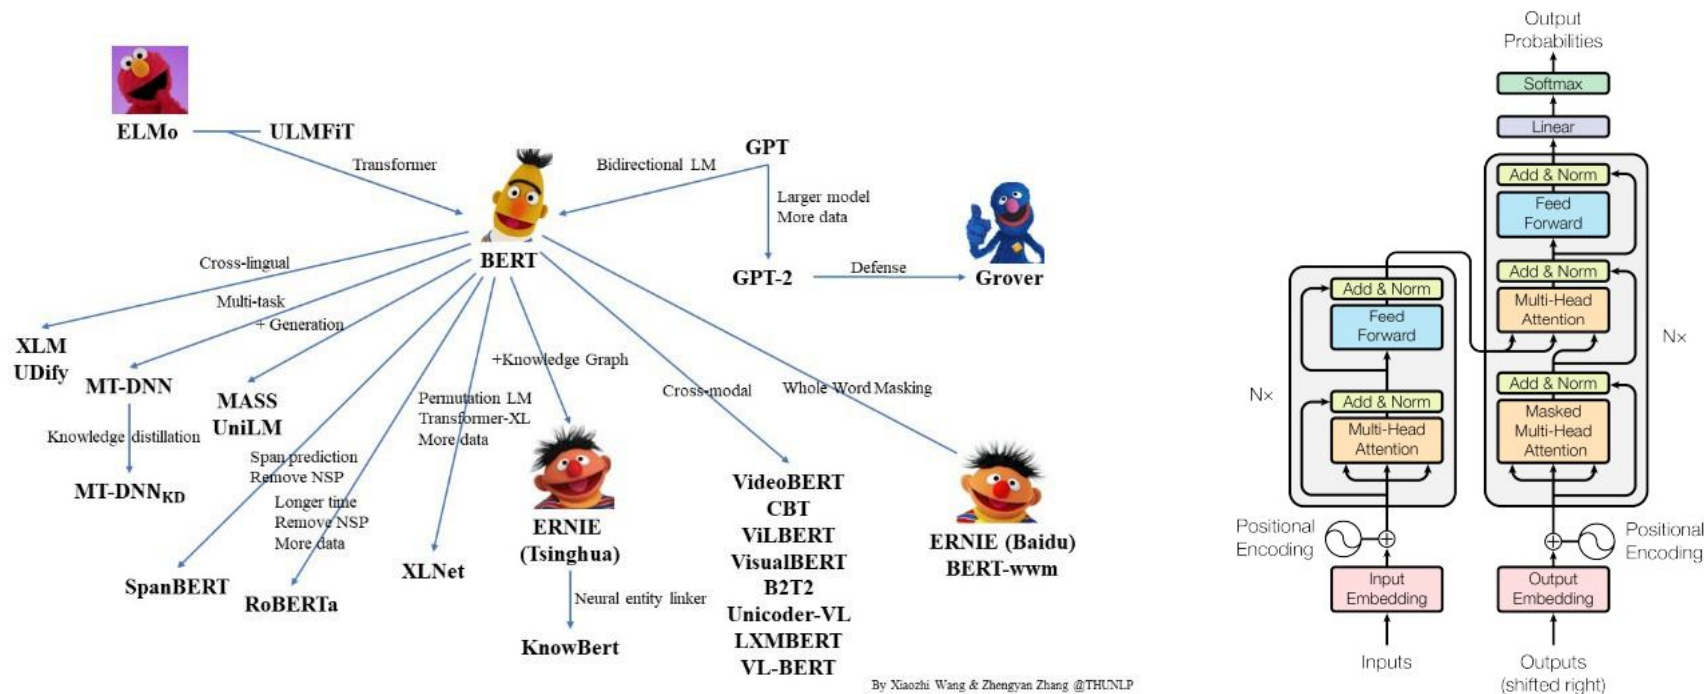



Figure 1: The Transformer - model architecture.

# Me : I want to be the cool kids

Try to run [XLM](#) model in GTX 1080, get out of memory errors instead

```
tokenizer = XLMTokenizer.from_pretrained('xlm-mlm-100-1280')
model = XLMWithLMHeadModel.from_pretrained('xlm-mlm-100-1280')
args.length = adjust_length_to_model(args.length, max_sequence_length=
model.to('cuda')
preprocessed_prompt_text = prepare_xlm_input(args, model, tokenizer,
encoded_prompt = tokenizer.encode(preprocessed_prompt_text, add_specia
encoded_prompt = encoded_prompt.to('cuda')
```



```
in linear
  output = input.matmul(weight.t())
RuntimeError: CUDA out of memory. Tried to allocate 1.6
8 GiB (GPU 0; 7.80 GiB total capacity; 2.17 GiB already
allocated; 1.08 GiB free; 2.21 GiB reserved in total b
y PyTorch)
(env) theblackcat102@lab:~/Documents/HW2$
```



# Why?

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V \quad O(n^2)$$



1. Self attention matrix grows quadratic with input length
2. Existing backpropagation strategy stores inputs and outputs of before each activation ( memory x 2 )



# Solution 1 : Sparse operation

Self attention between query and key vector is normally a sparse matrix. We can use sparse operation to save memory footprint.

But, sparse operation require ~~cool kids toys~~ fancy hardware support ( ie: TPU ) for fast performance; software implementation is slow as he\*\*



Cerebras WSE

1.2 Trillion transistors  
46,225 mm<sup>2</sup> silicon




Largest GPU

21.1 Billion transistors  
815 mm<sup>2</sup> silicon

# List of sparse transformer

- [Sparse Sinkhorn Attention](#) - google
- [Generating Long Sequences with Sparse Transformers](#) - openai
- [Adaptive Attention Span in Transformers](#) - facebook





## Solution 2: Recursive operation


A lite BERT (ALBERT) : recursive transformer layers and matrix factorized embedding to save memory.

But backpropagation still need to store inputs, outputs for each recursive iteration.

And the “real” culprit : self attention matrix size = length x length

# Trade computation for memory : LSH Attention

1. Share matrix for query and key ( follow up later )
2. Split self attention matrix into buckets , only calculate attention similar vectors




# But similarity calculation can be expensive

One way to do this is hashing :



Locality Sensitive Hashing!

As with the problem of hashing,  
there maybe collision





# Multiple random rotation hashing



Sphere Projected Points





Random Rotation 0



Random Rotation 1



Random Rotation 2



x: 0 2 1

y: 3 2 0

x: 0 2 1

y: 0 2 1

# Multiple random rotation hashing


suppose we want to split vector into  $b$  buckets, given a vector  $x$  we can find the hash code as follows

$R$  : random matrix with size of  $[d_k, b/2]$

$$h(x) = \operatorname{argmax}([xR; -xR])$$

we can use the max value of i-th index as bucket id

# Full illustration of LSH attention



# Query matrix = Key matrix


Number of queries and numbers of keys within a bucket maybe unequal.

Let Q matrix = K matrix and set key vector as the normalized version of query vector


$$h(k_j) = h(q_j) \text{ by setting } k_j = \frac{q_j}{\|q_j\|}$$

Leave out self attention, otherwise it will be largest value


Imbalance issues if use different matrix for query, key




(a) Normal



(b) Bucketed




(c) Q = K



(d) Chunked

# But LSH lose some “information” in each layers

Backpropagation updates for a given weight need its own input and output to compute.

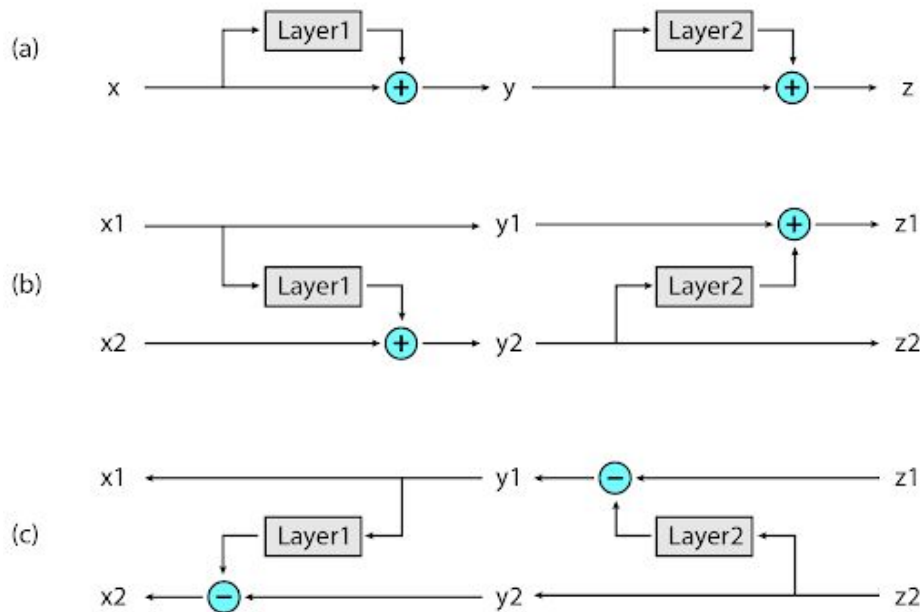




# Reversible Layers

RevNet tries to recover inputs from the next layers ( recover y from z )

We can make transformer reversible as well



$$Y_1 = X_1 + \text{Attention}(X_2)$$

$$Y_2 = X_2 + \text{FeedForward}(Y_1) \quad (9)$$

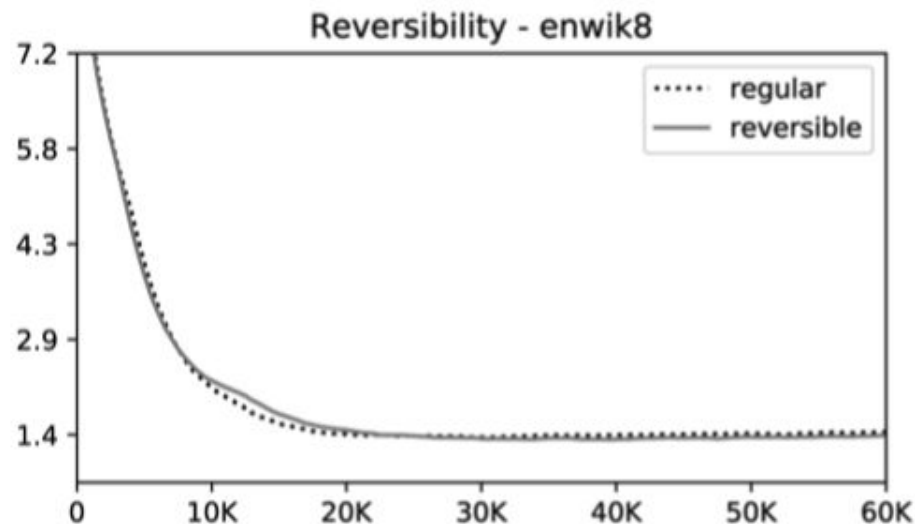
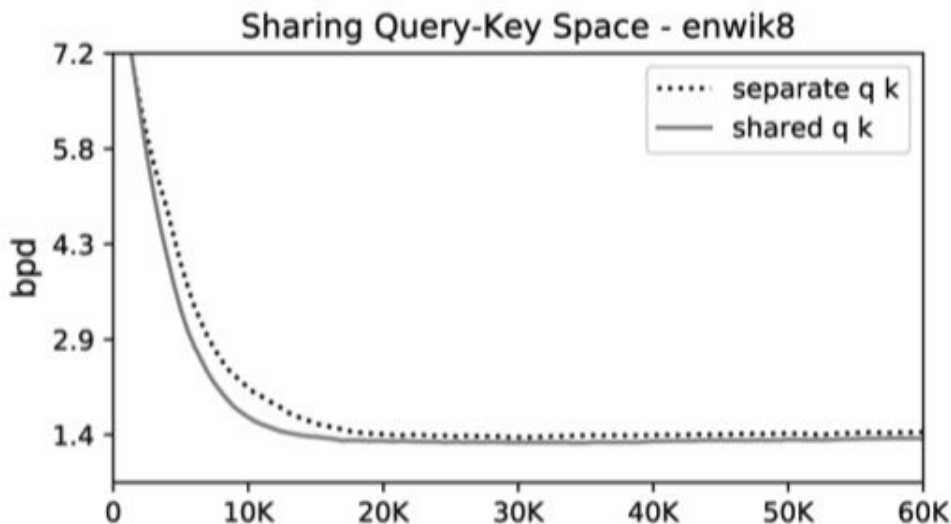
$$Y_2 = [Y_2^{(1)}; \dots; Y_2^{(c)}] = [X_2^{(1)} + \text{FeedForward}(Y_1^{(1)}); \dots; X_2^{(c)} + \text{FeedForward}(Y_1^{(c)})] \quad (10)$$

RevNet : The reversible residual network: Backpropagation without storing activations

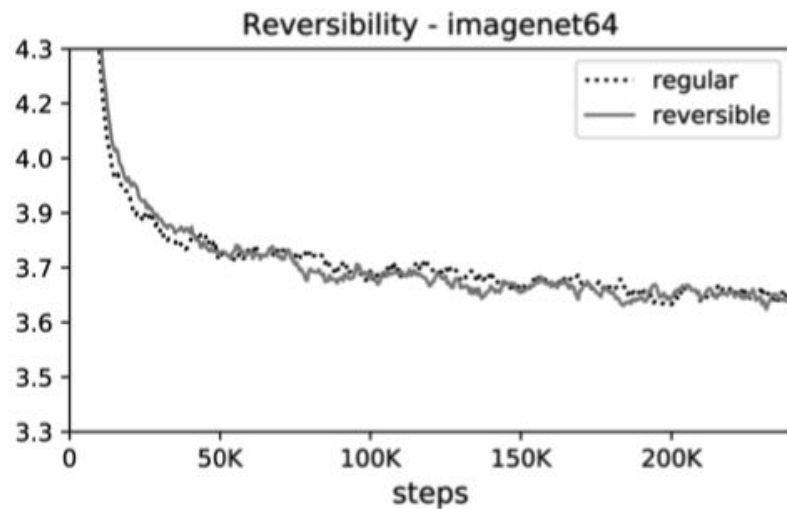

Implementation here: [https://github.com/lucidrains/reformer-pytorch/blob/master/reformer\\_pytorch/reversible.py](https://github.com/lucidrains/reformer-pytorch/blob/master/reformer_pytorch/reversible.py)

# Ablation study for Q=K, Reversibility

Bpd: bits per dim, basically negative log likelihood of outputs



# Ablation study for Q=K, Reversibility



# Memory, Time complexity

Table 3: Memory and time complexity of Transformer variants. We write  $d_{model}$  and  $d_{ff}$  for model depth and assume  $d_{ff} \geq d_{model}$ ;  $b$  stands for batch size,  $l$  for length,  $n_l$  for the number of layers. We assume  $n_c = l/32$  so  $4l/n_c = 128$  and we write  $c = 128^2$ .



Model Type	Memory Complexity	Time Complexity
Transformer	$\max(bld_{ff}, bn_h l^2)n_l$	$(bld_{ff} + bn_h l^2)n_l$
Reversible Transformer	$\max(bld_{ff}, bn_h l^2)$	$(bn_h ld_{ff} + bn_h l^2)n_l$
Chunked Reversible Transformer	$\max(bld_{model}, bn_h l^2)$	$(bn_h ld_{ff} + bn_h l^2)n_l$
LSH Transformer	$\max(bld_{ff}, bn_h l n_r c)n_l$	$(bld_{ff} + bn_h n_r l c)n_l$
Reformer	$\max(bld_{model}, bn_h l n_r c)$	$(bld_{ff} + bn_h n_r l c)n_l$

# Ablation study for number of hashing rounds

Table 2: Accuracies on the duplication task of a 1-layer Transformer model with full attention and with locality-sensitive hashing attention using different number of parallel hashes.

Train \ Eval					
	Full Attention	LSH-8	LSH-4	LSH-2	LSH-1
Full Attention	100%	94.8%	92.5%	76.9%	52.5%
LSH-4	0.8%	100%	99.9%	99.4%	91.9%
LSH-2	0.8%	100%	99.9%	98.1%	86.8%
LSH-1	0.8%	99.9%	99.6%	94.8%	77.9%

# Performance and speed scale vs input length, layers



# Take away

1. This paper use **hashing** to approximate attention matrix and introduce **reversible layers** to **save memory** at the expensive of additional computation
2. Pointed out in openreview, bucket size affects the model performance
3. Resource efficient ( less training time, run in CPU ) is a rising area in NLP as supervised natural language understanding performance close to near human level
  - a. [SuperGLUE : promote the development of effective, energy-efficient models for difficult NLU tasks.](#)
4. This paper only shows result from auto regressive training, other existing training method such as masking, masking in sequence-to-sequence learning need to evaluate as well to obtain the full view

# Questions

1. Why is shared query, key projection matrix required?
2. How do this paper reduce the self attention matrix complexity?



# References

- [\[2001.04451\] Reformer: The Efficient Transformer](#)
- [Reformer: The Efficient Transformer](#)
- [Reformer the efficient transformer - Youtube](#)
  
- Code implementation : [reformer-pytorch/reformer\\_pytorch at master · lucidrains/reformer-pytorch](#)